CS50's Introduction to Databases with SQL

OpenCourseWare

Donate (https://cs50.harvard.edu/donate)

Carter Zenke (https://carterzenke.me)
carter@cs50.harvard.edu
(https://github.com/carterzenke) (https://www.linkedin.com/in/carterzenke/)

David J. Malan (https://cs.harvard.edu/malan/)
malan@harvard.edu
(https://www.facebook.com/dmalan) (https://github.com/dmalan) (https://www.instagram.com/davidjmalan/) (https://www.linkedin.com/in/malan/) (https://www.reddit.com/user/davidjmalan) (https://www.threads.net/@davidjmalan) (https://twitter.com/davidjmalan)

Happy to Connect

Problem to Solve

You might recall from an earlier problem that LinkedIn (https://about.linkedin.com/) is "the world's largest professional network" with a mission to "connect the world's professionals to make them more productive and successful." Odds are, then, they use full-fledged database servers to ensure their platform is highly available (https://www.digitalocean.com/community/tutorials/what-is-high-availability) around the world.

In a file called `schema.sql` in a folder called `sentimental-connect`, write a set of SQL statements to design a MySQL database schema that LinkedIn could use.

Specification

Your task at hand is to create a MySQL database for LinkedIn from scratch. The implementation details are up to you, though you should minimally ensure that your database meets the platform's specification and that it can represent the given sample data. You're welcome to use, or improve upon, your design of a SQLite database—just keep in mind that you'll now have a new set of types (https://dev.mysql.com/doc/refman/8.0/en/data-types.html) at your disposal!

Platform

Users

The heart of LinkedIn's platform is its people. Your database should be able to represent the following information about LinkedIn's users:

- Their first and last name
- Their username
- Their password

Keep in mind that, if a company is following best practices, application passwords are "hashed (https://en.wikipedia.org/wiki/Hash_function)." No need to worry about hashing passwords here, though it might be helpful to know that some hashing algorithms can produce strings up to 128 characters long.

Schools and Universities

LinkedIn also allows for official school or university accounts, such as that for Harvard (https://www.linkedin.com/school/harvard-university/), so alumni (i.e., those who've attended) can identify their affiliation. Ensure that LinkedIn's database can store the following information about each school:

- The name of the school
- The type of school
- The school's location
- The year in which the school was founded

You should assume that LinkedIn only allows schools to choose one of three types: "Primary," "Secondary," and "Higher Education."

Companies

LinkedIn allows companies to create their own pages, like the one for LinkedIn itself (https://www.linkedin.com/company/linkedin/), so employees can identify their past or current employment with the company. Ensure that LinkedIn's database can store the following information for each company:

- The name of the company
- The company's industry
- The company's location

You should assume that LinkedIn only allows companies to choose from one of three industries: "Technology," "Education," and "Business."

Connections

And finally, the essence of LinkedIn is its ability to facilitate connections between people. Ensure LinkedIn's database can support each of the following connections.

▼ Connections with People

LinkedIn's database should be able to represent mutual (reciprocal, two-way) connections between users. No need to worry about one-way connections user A "following" user B without user B "following" user A.

▼ Connections with Schools

A user should be able to create an affiliation with a given school. And similarly, that school should be able to find its alumni. Additionally, allow a user to define:

- The start date of their affiliation (i.e., when they started to attend the school)
- The end date of their affiliation (i.e., when they graduated), if applicable
- The type of degree earned/pursued (e.g., "BA", "MA", "PhD", etc.)

▼ Connections with Companies

A user should be able to create an affiliation with a given company. And similarly, a company should be able to find its current and past employees. Additionally, allow a user to define:

- The start date of their affiliation (i.e., the date they began work with the company)
- The end date of their affiliation (i.e., when left the company), if applicable

Sample Data

Your database should be able to represent...

- A user, **Claudine Gay (https://en.wikipedia.org/wiki/Claudine_Gay)**, whose username is

"claudine" and password is "password".

- A user, **Reid Hoffman (https://en.wikipedia.org/wiki/Reid_Hoffman)** whose username is "reid" and password is "password".
- A school, **Harvard University**, which is a university located in Cambridge, Massachusetts, founded in 1636.
- A company, **LinkedIn**, which is a technology company headquartered in Sunnyvale, California.
- Claudine Gay's **connection with Harvard**, pursuing a PhD from January 1st, 1993, to December 31st, 1998.
- Reid Hoffman's **connection with LinkedIn**, with title "CEO and Chairman", from January 1st, 2003 to February 1st, 2007

Advice

- Consider the full range of MySQL's supported types, which are documented in the MySQL 8.0 reference manual at dev.mysql.com/doc/refman/8.0/en/data-types.html (https://dev.mysql.com/doc/refman/8.0/en/data-types.html).
- Also consider the reference manual's advice on choosing the right type for a column, documented at dev.mysql.com/doc/refman/8.0/en/choosing-types.html (https://dev.mysql.com/doc/refman/8.0/en/choosing-types.html).
    - Among the high-level pieces of advice is to choose the most precise type for your use case. For instance, if you know an integer column will store only positive values, you should consider modifying the integer type with `UNSIGNED` (e.g., `INT UNSIGNED` or `TINYINT UNSIGNED`) to get the most range out of your type.

Usage

To use `mysql`, you first need to start a MySQL server, as with:

```
docker container run --name mysql -p 3306:3306 -v /workspaces/$RepositoryName:/mn
```

You can then connect to the server with:

```
mysql -h 127.0.0.1 -P 3306 -u root -p
```

Type **crimson** as your password.

If this is your first time logging into your `mysql` server, you'll need to create a new database on the server. At your `mysql>` prompt, try the following:

```
CREATE DATABASE `linkedin`;
```

You should see something along the lines of "Query OK, 1 row affected." Afterwards, to ensure your future SQL statements are run on your new `linkedin` database, execute the following:

```
USE `linkedin`;
```

Now recall other MySQL statements from lecture which might help you navigate! You can always type `quit` to close your connection to the MySQL database.

How to Test

While `check50` exists for this problem, only you can ensure your database meets the platform's specification and that it can store the sample data efficiently. Consider whether your database is fully normalized!

Correctness

```
check50 cs50/problems/2023/sql/sentimental/connect
```

How to Submit

In your terminal, execute the below to submit your work.

```
submit50 cs50/problems/2023/sql/sentimental/connect
```